

LifeLink AI: A Trustworthy Hybrid Clinical Decision-Support Architecture for Real-Time Emergency Blood Matching in Digital and Precise Medicine

Shivang Mishra, Lakshya Sahu, Kushagra Bhargava, Anshul, Archisha Nigam, Mokshi Jain

VIT Bhopal University, Bhopal, India

Abstract—Emergency blood provision during acute haemorrhagic trauma, perioperative crises, and obstetric haemorrhage demands sub-minute coordination under uncompromising clinical safety constraints. Traditional blood procurement across developing healthcare ecosystems relies on fragmented telephonic communication, static web directories, and unverified broadcast appeals, introducing dispatch latencies of 20 to 60 minutes that substantially elevate mortality during the critical ‘golden hour.’ This paper presents **LifeLink AI**, a trustworthy hybrid clinical decision-support architecture engineered for real-time emergency blood matching within the paradigm of Digital and Precise Medicine. The architectural foundation is a formally defined **two-layer safety contract** that enforces a categorical boundary between deterministic medical invariants and probabilistic artificial intelligence. Layer 1 executes an immunohematological verification gate implementing a pure-Python $O(1)$ ABO/Rh compatibility mapping and statutory health criteria—a 56-day whole-blood recovery cooldown, minimum body mass ≥ 45 kg, and verified donor eligibility—that are computationally isolated from and cannot be overridden by any machine-learning component. Layer 2 applies a calibrated logistic-regression donor-response propensity model combined with PostGIS geospatial proximity decay and multi-criteria composite scoring to generate an advisory ranking $S \in [0, 1]$. The propensity model (StandardScaler + balanced LogisticRegression) is trained on the UCI Blood Transfusion Service Center benchmark (748 records, 5-fold stratified cross-validation), achieving ROC-AUC = 0.7521, Recall = 0.7528, Precision = 0.3907, F1 = 0.5144, and Brier Score = 0.2055. The platform is containerised across six microservices (Next.js 14, FastAPI, scikit-learn inference service, PostgreSQL 15 + PostGIS 3.3, Redis 7, Nginx) and validated across 56 automated integration tests with a 100% pass rate. Benchmarking demonstrates mean matching pipeline latencies of 47.2 ms on the AI path and 12.4 ms on the deterministic fallback path. LifeLink AI demonstrates how trustworthy hybrid AI architectures can bridge critical logistics bottlenecks in emergency precision healthcare while maintaining non-negotiable patient safety guarantees.

Keywords—Digital and precise medicine; trustworthy AI; clinical decision support; emergency blood matching; immunohematology gate; donor response propensity; hybrid architecture; microservices; explainable AI; logistic regression; RFM model; PostGIS.

I. INTRODUCTION

Haemorrhagic shock arising from traumatic injury, postpartum haemorrhage, major cardiovascular surgery, and acute oncological cytopenias represents one of the foremost causes of preventable mortality in emergency medicine worldwide [1]. Clinical evidence across trauma resuscitation consistently demonstrates that patient survival probability decays non-linearly as time-to-transfusion increases; within the critical ‘golden hour,’ every additional ten-minute delay in procuring cross-match-compatible blood units correlates directly with exponential increases in multiorgan failure risk and mortality [2], [3]. Globally, the World Health Organization (WHO) estimates that approximately 118.5 million blood donations are collected annually, yet profound and often lethal disparities persist in the logistical distribution and emergency allocation of blood products across regional healthcare networks [4].

In India, where annual blood demand approaches 14 million units against an estimated collection of 12 to 15 million units [5], the clinical challenge is rarely an absolute national shortage; rather, it is a severe structural and temporal coordination failure caused by geographic maldistribution, cold-chain fragmentation, and pervasive

information asymmetry across regional healthcare providers. Emergency blood procurement across municipal healthcare networks currently suffers from three pervasive, compounding structural defects.

- **Manual and Fragmented Telephonic Coordination:** Hospital emergency departments and patient relatives frequently resort to uncoordinated telephone cascades across regional blood banks and voluntary donor lists. This ad-hoc process consumes 20 to 60 minutes per requisition and lacks real-time verification of whether a target facility holds unreserved, non-expired compatible units in cold storage.
- **Static Directories and Broadcast Fatigue:** Web-based directories list contact details but provide no real-time availability data, leading to high rejection rates from exhausted inventories. Mass broadcast appeals via social media and messaging platforms foster donor fatigue, progressively desensitising eligible volunteers to genuine life-threatening emergencies.
- **Absence of Intelligent Predictive Stratification:** Existing hospital information systems treat all eligible donors identically, ignoring historical donation recency, frequency, and geographic transit times. Consequently, dispatch coordinators waste vital minutes contacting

registered individuals who are statistically unlikely to respond or travel to the clinical facility within the therapeutic window.

Addressing these bottlenecks rigorously requires an engineering approach grounded in **Digital and Precise Medicine**—the application of algorithmic and computational precision to the logistics of life-saving biological products under acute temporal constraints. However, introducing artificial intelligence into emergency medicine creates substantial clinical and medicolegal risks. Machine learning models deployed in high-stakes clinical workflows are susceptible to distribution shifts, probabilistic overconfidence, and catastrophic failure under rare-event conditions [6], [7]. In blood transfusion specifically, an algorithmic error recommending an ABO-incompatible unit can induce acute intravascular haemolysis, acute kidney injury, systemic inflammatory shock, and death within minutes of infusion. No probabilistic model—regardless of its training performance—can be permitted to govern this immunohematological determination.

To resolve this fundamental tension between algorithmic efficiency and inviolable patient safety, this paper presents **LifeLink AI**, a trustworthy hybrid clinical decision-support system. LifeLink AI implements a **two-layer safety contract** that enforces an absolute, computationally enforced boundary: biological compatibility and statutory eligibility are governed exclusively by deterministic medical logic, while machine learning operates entirely as an advisory donor-response propensity layer that influences only operational ranking—never medical decisions.

Primary Research Contributions:

- *The Two-Layer Safety Contract*: A formalized architectural design pattern that mathematically isolates zero-tolerance deterministic medical invariants (ABO/Rh compatibility, statutory recovery cooldowns, weight gates) from probabilistic AI components across containerized network boundaries, with no shared code paths.
- *Hybrid Multi-Criteria Composite Scoring*: A calibrated scoring formulation integrating deterministic compatibility metrics, PostGIS geospatial proximity decay, operational facility availability, and machine-learning donor response propensity into an auditable ranking function with explicit, interpretable weight assignments.
- *Calibrated ML Propensity Modeling*: Implementation and empirical evaluation of a class-balanced logistic-regression response model trained on the UCI Blood Transfusion benchmark, demonstrating how

clinical sensitivity (recall = 0.7528) is appropriately prioritised over raw accuracy in asymmetric high-stakes classification.

- *Production-Grade Microservice Topology*: A fully containerized 6-tier architecture validated by 56 automated integration tests (100% pass rate), sub-100 ms execution latencies, and pessimistic concurrency locking (SELECT ... FOR UPDATE) preventing double-allocation of finite cold-chain inventory under concurrent emergency load.
- *Relationship-Based Access Control (ReBAC)*: A 7-role multi-tenant governance model ensuring secure, authenticated interaction across donors, hospitals, blood banks, and administrative authorities, with complete immutable audit logging of every matching run and inventory transaction.

II. RELATED WORK AND LITERATURE REVIEW

A. Blood Supply Chain Logistics and Management

The management of human blood products has been extensively studied as a perishable inventory routing problem exhibiting severe time-temperature constraints, stochastic demand, and heterogeneous supply points. Stanger et al. [8] surveyed inventory management strategies across European transfusion services, identifying that blood wastage (outdating) and acute supply shortages stem primarily from fragmented inventory visibility and non-integrated ordering systems across regional hospital networks. Early technological interventions focused on computerized physician order entry (CPOE) and computerized blood ordering systems (CBOS). Heitmiller et al. [9] demonstrated that algorithmic clinical guidelines embedded within CPOE software reduced unnecessary cross-matches by 42% and significantly shortened procurement turnaround times in elective surgical settings—a finding that motivated the deterministic gate design in LifeLink AI.

In recent years, Internet of Things (IoT) sensor networks and RFID tagging have been introduced to automate cold-chain tracking, temperature excursion monitoring, and barcode-driven chain-of-custody [10]. However, these systems address predominantly static facility inventory management and supply-chain traceability rather than dynamic, real-time multi-facility emergency coordination across voluntary donor populations during acute trauma events. In India, government-backed platforms such as *e-Rakt Kosh* [11] provide centralized web-based directories of licensed blood banks; however, they function primarily as passive administrative registries and lack automated spatial

candidate matching, real-time donor availability verification, and predictive response modeling—precisely the capabilities that LifeLink AI contributes.

B. Machine Learning in Donor Recruitment and Retention

Predicting donor lapse, return probability, and emergency response behavior has engaged the data science and transfusion medicine communities since the foundational publication of the Recency, Frequency, Monetary (RFM) framework for blood transfusion by Yeh et al. [12]. Their empirical study on 748 Blood Transfusion Service Center (BTSC) records in Hsin-Chu City, Taiwan established that recent and frequent donors exhibit measurably higher future response propensity, yielding a population-level positive class prevalence of 23.8%—a pronounced class imbalance that fundamentally motivates the balanced class-weight strategy adopted in LifeLink AI.

Gilliss et al. [13] evaluated artificial neural networks and logistic regression for predicting first-time donor return rates in voluntary donation programs, identifying that communication channel timeliness and operational convenience significantly influence donor loyalty. Subsequent meta-analyses of classifier performance across donor retention datasets [14] reveal that while complex non-linear ensemble models occasionally yield marginal gains in area under the ROC curve, regularized logistic regression consistently provides superior probability calibration, monotonic score stability under feature perturbation, and model-level intrinsic interpretability when operating on low-dimensional behavioral feature sets—a critical advantage in safety-critical clinical applications where model transparency is non-negotiable. Ramachandran et al. [15] corroborated this finding, observing that unconstrained deep neural networks exhibit catastrophically unreliable out-of-distribution behavior in high-stakes medical logistics tasks, recommending calibrated linear models for domains where worst-case failure has irreversible consequences.

C. Trustworthy and Explainable AI in Digital Healthcare

The translation of artificial intelligence from benchmark research into frontline clinical practice is severely constrained by the ‘black box’ problem. In high-stakes medicine, clinicians, transfusion coordinators, and hospital ethics committees cannot accept opaque recommendations from deep neural networks without verifiable causal justifications [6], [16]. Doshi-Velez and Kim [17] articulated a rigorous taxonomy of explainability distinguishing post-hoc interpretability (e.g., LIME or SHAP approximations of complex models) from intrinsic interpretability, wherein the model’s structural equations are directly comprehensible to domain experts without requiring secondary approximation methods that may themselves introduce distortions. LifeLink AI adopts intrinsic interpretability through logistic regression: every scoring weight and propensity probability is directly readable from the model coefficient vector.

The European Union High-Level Expert Group on AI [18] and the IEEE Global Initiative on Ethics of Autonomous and Intelligent Systems [19] have formulated overlapping core tenets for Trustworthy AI in high-stakes domains: clinical validity, non-maleficence, transparency, accountability, robustness, and privacy protection. The European AI Act [20] further categorizes medical AI systems with direct patient impact as high-risk systems requiring mandatory human oversight, conformity assessment, and logging of decisions. LifeLink AI operationalizes these principles through: (a) architecturally enforced determinism for biological safety decisions; (b) a three-level resilient fallback hierarchy that preserves operation under partial AI infrastructure failure; (c) immutable audit logs of every matching run; and (d) human-in-the-loop dispatch, where clinicians and coordinators retain final dispatch authority [21], [22].

D. Comparative Positioning

Table I situates LifeLink AI against representative prior systems across six critical architectural dimensions. Static directory platforms (e.g., e-Rakt Kosh, RaktDaan) provide no algorithmic matching or predictive stratification. Pure ML platforms that apply probabilistic models directly to compatibility decisions introduce unacceptable patient risk—a failure mode documented in early clinical AI deployments [7]. LifeLink AI uniquely combines the deterministic safety guarantees of rule-based systems with the operational optimization capabilities of machine learning, while additionally providing production-grade concurrency safety, multi-tenant governance, and a transparent three-level fallback hierarchy absent from all surveyed prior systems.

TABLE I—COMPARATIVE TAXONOMY OF BLOOD COORDINATION SYSTEMS

Dimension	Static Registries	Pure ML Platforms	LifeLink AI (Proposed)
Blood Compatibility	Manual / External	Probabilistic (ML)	Deterministic O(1) Gate
Donor Prioritization	FIFO / Unranked	Black-Box Ensemble	Calibrated Linear Scoring
Explainability	N/A	Post-Hoc (SHAP/LIME)	Intrinsic (Coefficients)
Fallback Behavior	Manual Intervention	Unhandled Exception	3-Level Resilient Chain
Concurrency Safety	None (Siloed)	None	SELECT FOR UPDATE Lock
Statutory Health Gate	Self-Declared	Omitted / Input Feature	Strict Pre-Filter Layer
Audit Trail	None	None / Informal	Immutable Ledger

III. PROBLEM STATEMENT AND CLINICAL MOTIVATION

In severe trauma resuscitation, transfusion delays compound systemic coagulopathy and metabolic acidosis, creating the lethal triad of hypothermia, acidosis, and coagulopathy that rapidly becomes irreversible without rapid blood product infusion. When a clinical emergency blood requisition is initiated, two distinct biological and logistical entity classes can fulfill the demand: (1) *Cold-Chain Blood Bank Inventory*: preserved units of Whole Blood, Packed Red Blood Cells (PRBC), Fresh Frozen Plasma (FFP), or Platelets maintained in licensed blood banks; and (2) *Voluntary Citizen Donors*: registered eligible individuals who must travel to a designated collection center for emergency phlebotomy.

The core algorithmic challenge is to *simultaneously* discover, filter, and prioritize both candidate pools within seconds under the following non-negotiable constraints: (a)

Zero medical risk—biological compatibility must be determined with absolute certainty through deterministic immunohematological rules; (b) **Minimal transit time**—geographic proximity must be quantified and weighted to maximize the probability of receiving the unit within the therapeutic window; (c) **Maximum response certainty**—donors most likely to respond must be surfaced first to avoid wasting contact time; and (d) **Concurrency safety**—finite blood inventory must never be double-allocated to competing simultaneous emergency requests from different hospitals.

The secondary challenge is architectural trustworthiness: the system must remain operable under AI infrastructure failure, must log every decision with immutable traceability, must enforce multi-tenant security preventing cross-facility data leakage, and must be explainable to attending physicians who require transparent justification for any algorithmic recommendation they act upon in a life-or-death clinical scenario.

IV. SYSTEM MODEL AND MATHEMATICAL FORMULATION

A. Emergency Request Formulation

Let an emergency blood requisition be formally defined as a typed tuple:

$$E = (req_{id}, \beta_{rec}, \gamma_{comp}, u_{req}, \lambda_{dest}, \omega_{urg}, t_0) \quad (1)$$

where $req_{id} \in \text{UUID}$ is a platform-unique cryptographic identifier; $\beta_{rec} \in \blacksquare$ denotes the recipient ABO/Rh blood group drawn from the universal set of eight clinically recognized groups:

$$\blacksquare = \{ O-, O+, A-, A+, B-, B+, AB-, AB+ \} \quad (2)$$

$\gamma_{comp} \in \{\text{WHOLE_BLOOD, RBC, PLASMA, PLATELETS}\}$ denotes the clinician-prescribed blood component fraction; $u_{req} \in [1, \uparrow 20]$ is the requested volume in units; $\lambda_{dest} = (\phi_{dest}, \psi_{dest})$ denotes the WGS-84 geospatial coordinates of the receiving facility; $\omega_{urg} \in \{\text{LOW, MEDIUM, HIGH, CRITICAL}\}$ defines the clinical urgency tier; and t_0 is the UTC submission timestamp.

B. Candidate Search Space

The candidate search space comprises two disjoint pools. Let $\Omega_D = \{d_1, \dots, d_N\}$ denote registered voluntary donors and $\Omega_K = \{k_1, \dots, k_M\}$ denote licensed blood bank facilities. Each donor profile is parameterized by:

$$d_i = (id_i, \beta_i, w_i, a_i, e_i, \lambda_i, t_{last}, F_i, T_i) \quad (3)$$

where w_i (kg) is body mass; $a_i, e_i \in \{0, 1\}$ are real-time availability and medical eligibility flags; t_{last} is the UTC timestamp of the most recent whole-blood donation; F_i is lifetime donation count; and T_i is months elapsed since first

registration. Each blood bank candidate is parameterized by:

$$k_j = (id_j, name_j, lic_j, \lambda_j, v_j, I_j) \quad (4)$$

where I_j is the cold-chain inventory ledger associating component and blood group to available and reserved unit counts. Geospatial proximity is quantified by the Haversine great-circle metric $d_{hav}(\lambda_1, \lambda_2)$, with active spatial search radius $R_{search} \in \{15, 25, 50, 100\}$ km executed via PostGIS ST_DWithin with a GiST-indexed geometry column.

TABLE II—MATHEMATICAL NOTATION AND VARIABLE DEFINITIONS

Symbol	Domain	Clinical / Architectural Meaning
E	Tuple	Emergency blood requisition entity
$\beta_{■■■■}, \beta_{■■■■}$	■	Recipient and donor ABO/Rh blood groups
$\gamma_{■■■■}$	Component Set	Prescribed fraction (Whole Blood, RBC, Plasma, Platelets)
$u_{■■■■}$	[1, 20]	Blood volume in units requested by attending physician
$\lambda = (\phi, \psi)$	WGS-84	Geospatial coordinates (latitude, longitude)
$d_{■■■■}(\lambda_{■■}, \lambda_{■■})$	km	Haversine great-circle distance between two locations
$R_{■■■■■■}$	{15,25,50,100} km	Active spatial PostGIS search radius
$\Delta t_{■■■■}$	≥ 56 days	Recovery cooldown since last whole-blood donation
$P_{■■}$	[0.0, 1.0]	Calibrated ML donor response propensity probability
$P_{■■■■}$	[0.0, 1.0]	Linear spatial proximity decay over active radius
$S_{■■■■■■}$	[0.0, 1.0]	Multi-criteria composite donor ranking score
$S_{■■}$	[0.0, 1.0]	Multi-criteria composite blood bank ranking score
$\alpha_{■■}, \alpha_{■■}$	■■	Inverse class frequency weights for balanced logistic loss

V. PROPOSED LIFELINK AI ARCHITECTURE




Fig. 1. LifeLink AI two-layer safety contract: categorical separation between deterministic immunohematological pre-filtering (Layer 1) and advisory AI propensity scoring (Layer 2). Medical decisions are exclusively Layer 1 outputs; Layer 2 influences operational ranking only.

A. Microservice Topology and Container Orchestration

LifeLink AI is engineered as an enterprise-grade distributed microservice platform composed of six isolated Docker containers, each with a strictly defined responsibility boundary and network exposure policy: (1) **lifelink-nginx** (port 80)—Nginx reverse proxy providing SSL termination, request routing, and rate limiting; (2) **lifelink-frontend** (port 3000)—Next.js 14.2.5 App Router single-page application providing role-specific user interfaces for donors, hospital administrators, blood bank managers, and super-administrators; (3) **lifelink-backend** (port 8000)—FastAPI 0.111.1 orchestration engine on Python 3.11 implementing all business logic, matching pipeline coordination, and database transactions; (4) **lifelink-ai-service** (port 8001)—a dedicated scikit-learn 1.5.1 inference microservice executing the calibrated logistic regression model in an isolated process context; (5) **lifelink-postgres** (port 5432)—PostgreSQL 15 with PostGIS 3.3 spatial extension serving as the system of record; and (6) **lifelink-redis** (port 6379)—Redis 7 serving as session cache and API rate-limiter. All inter-service communication runs on an isolated Docker bridge network; only Nginx exposes a public port.

B. Relational Schema and Spatial Data Engineering

The database architecture comprises 11 relational tables managed under migration control: *users*, *user_roles*, *donors*, *hospitals*, *blood_banks*, *blood_inventory*, *inventory_history*, *emergency_requests*, *match_runs*, *match_candidates*, and *donor_emergency_responses*. Spatial attributes (location_geom) are stored as native PostGIS geometry objects with SRID=4326 (WGS-84) and indexed using Generalized Search Trees (GiST), enabling $O(\log N)$ spatial radius queries via ST_DWithin. Database-level CHECK constraints enforce physiological bounds ($weight_kg \geq 45$, $units_required$ between 1 and 20). The *inventory_history* table operates as an immutable append-only audit ledger—UPDATE and DELETE operations are disallowed at the database role level—guaranteeing complete traceability for every blood unit reserved, released, or dispatched.

C. Relationship-Based Access Control (ReBAC)

The system enforces a granular 7-role authorization hierarchy implemented at three independent enforcement layers: (1) JWT cryptographic signature verification (HMAC-SHA256, 30-minute access tokens, 7-day rolling refresh tokens); (2) FastAPI endpoint dependency injection validating role membership and active status; and (3) SQLAlchemy query filters scoping all data access to the requesting user's institutional tenancy identifier. The role hierarchy is: `SUPER_ADMIN` $\rightarrow$ `ADMIN` $\rightarrow$ `{HOSPITAL_ADMIN, BLOOD_BANK_MANAGER}` $\rightarrow$ `{HOSPITAL_STAFF, BLOOD_BANK_STAFF}` $\rightarrow$ `DONOR`. Cross-tenant data access is rejected at HTTP layer with 403 Forbidden before any database query is executed.

VI. DETERMINISTIC MEDICAL COMPATIBILITY AND ELIGIBILITY LAYER

Immunohematological compatibility is an immutable biological law, not a probabilistic pattern. Administering ABO-incompatible blood triggers complement-mediated intravascular haemolysis, which can be fatal within minutes. Accordingly, LifeLink AI implements the deterministic compatibility gate as a pure-Python, machine-learning-free module (*backend/app/core/medical.py*) with zero external dependencies and $O(1)$ worst-case time complexity. For Whole Blood and Packed Red Blood Cells (PRBC), compatible donor blood groups for a given recipient type β_{rec} are defined by the set:

$$C_{RBC}(\beta_{rec}) = \{ \beta_{don} \in \blacksquare \blacksquare \text{ donor RBC antigenically compatible with } \beta_{rec} \} \quad (5)$$

Implemented as a static Python dictionary keyed by recipient group, each entry maps to the frozenset of valid donor groups. For Fresh Frozen Plasma (FFP), where the dominant compatibility constraint inverts to donor antibody profile rather than antigen surface, the mapping is separately defined:

$$C_{Plasma}(\beta_{rec}) = \{ \beta_{don} \in \blacksquare \blacksquare \text{ donor plasma compatible with } \beta_{rec} \} \quad (6)$$

Under FFP rules, AB plasma is universally compatible (universal donor), while O recipients can only receive O plasma. A candidate donor d_i passes the deterministic eligibility gate if and only if all five of the following conditions are simultaneously satisfied:

$$Gate(d_i) = [\beta_i \in C(\beta_{rec})] \wedge [e_i = 1] \wedge [a_i = 1] \wedge [\Delta t_{cool} \geq 56] \wedge [w_i \geq 45] \quad (7)$$

The 56-day whole-blood recovery period is a WHO-mandated statutory requirement; the 45 kg minimum body mass is enforced both at the application layer and at the database level via a CHECK constraint. Donors failing

any single predicate are unconditionally excluded from the candidate pool before any scoring computation is performed. This design ensures that the machine-learning propensity layer is structurally incapable of influencing biological eligibility decisions.

VII. AI DONOR RESPONSE PROPENSITY MODEL

A. Precise Role and Clinical Boundaries of AI

The artificial intelligence component in LifeLink AI is explicitly engineered and documented as an **advisory response propensity estimator** operating within Layer 2. Its singular, clearly bounded function is to estimate the conditional probability that a registered voluntary donor, having already passed the deterministic Layer 1 gate, will accept an emergency outreach alert and travel to the designated collection facility within the therapeutic window. This scope definition is not merely a policy statement; it is enforced through system architecture: the AI service receives only behavioral features (R, F, T) as input and returns only a probability in $[0, 1]$ as output. It does *not* receive blood group data, patient identity, or eligibility flags; it cannot modify the candidate pool; it cannot initiate or suppress any dispatch action.

B. Dataset Provenance and Feature Engineering

The propensity model is trained on the UCI Blood Transfusion Service Center (BTSC) benchmark [12] (OpenML dataset ID 1464, licensed CC BY 4.0). The dataset contains $N = 748$ anonymized donor records from the Hsin-Chu City, Taiwan blood transfusion service, collected over a period spanning 1988 to 1994. The target binary variable $y \in \{0, 1\}$ indicates whether the donor donated blood in March 2007, yielding a class distribution of $N_{pos} = 178$ (23.8%) positive instances and $N_{neg} = 570$ (76.2%) negative instances—a significant 3.2:1 class imbalance.

The original dataset provides four features: Recency (R, months since last donation), Frequency (F, total lifetime donations), Monetary (M, total blood volume donated in mL as $M = 250 \times F$), and Time (T, months since first donation). Preliminary statistical analysis revealed that Monetary exhibits exact linear dependence on Frequency with Pearson correlation coefficient $r = 1.000$, confirming perfect collinearity. Retaining M in the feature vector would introduce a singular Gram matrix and inflate coefficient variance without adding predictive information. Accordingly, Monetary was pruned, yielding the final three-dimensional feature vector:

$$x = [R, F, T]^T \in \blacksquare^3 \quad (8)$$

C. Model Architecture and Balanced Loss Formulation

Raw features x are standardized via a fitted StandardScaler to zero mean and unit variance, producing standardized features:

$$x_{\text{std}} = \Sigma^{-1/2} (x - \mu) \quad (9)$$

where $\mu \in \mathbb{R}^3$ is the empirical feature mean vector and $\Sigma^{1/2}$ is the diagonal feature standard deviation matrix. Donor response propensity is then modeled by binary logistic regression:

$$P(y=1 | x) = \sigma(w^T x_{\text{std}} + b) = 1 / (1 + e^{-(w^T x_{\text{std}} + b)}) \quad (10)$$

where $w \in \mathbb{R}^3$ is the learned weight vector and $b \in \mathbb{R}$ is the bias term. To address the structural class imbalance (23.8% positive rate), the empirical risk minimization objective incorporates inverse class frequency weights:

$$L(w, b) = -1/N \sum_i [\alpha_{\text{pos}} y_i \log p_i + \alpha_{\text{neg}} (1 - y_i) \log(1 - p_i)] + \lambda/2 \|w\|_2^2 \quad (11)$$

$$\alpha_{\text{pos}} = N / (2 N_{\text{pos}}) = 2.101; \alpha_{\text{neg}} = N / (2 N_{\text{neg}}) = 0.656 \quad (12)$$

where $\lambda = 1/C$ with regularization strength $C = 1.0$ (the sklearn default), providing L2 regularization to prevent overfitting. The $\alpha_{\text{neg}} > 1$ weight penalizes False Negatives (missed willing donors) more heavily than False Positives (spurious alerts), directly operationalizing the clinical priority of donor recall over contact precision. Parameter optimization is performed using the L-BFGS quasi-Newton algorithm, which is well-suited to small, dense feature matrices and guarantees convergence to the global optimum of the convex log-loss objective.

D. Three-Level Resilient Fallback Hierarchy

To guarantee uninterrupted emergency operations under AI infrastructure failures, LifeLink AI implements a three-level degradation hierarchy with no single points of failure in the matching pipeline:

Level 1 (Primary AI Service Healthy): The FastAPI backend forwards donor RFT features via HTTP POST to the AI inference microservice (port 8001), which applies the fitted StandardScaler and LogisticRegression to return a calibrated propensity probability $P_{\text{ML}} \in [0, 1]$. This path achieves a mean latency of 47.2 ms.

Level 2 (Model I/O Failure, Microservice Reachable): If the inference call fails due to a model artifact loading error or internal service exception, the AI microservice activates an in-process closed-form RFM heuristic:

$$P_{\text{fb}} = \min(0.95, \max(0.05, 0.40 + \max(0, 0.35 - 0.01R) + \min(0.25, 0.05F))) \quad (13)$$

Level 3 (Microservice HTTP Unreachable, 503 Response): If the AI service is completely unavailable, the

backend assigns a neutral prior $P_{\text{ML}} = 0.50$ and records `model_version = 'deterministic-fallback'` in the audit record. The matching pipeline continues executing with the reduced score:

$$S_{\text{geo}} = 0.40 \times C + 0.30 \times P_{\text{geo}} + 0.20 \times A + 0.05 \quad (14)$$

This three-level hierarchy ensures that even total AI service failure degrades gracefully to deterministic-only matching rather than producing an unhandled error or halting the emergency dispatch pipeline.

VIII. MATCHING AND RANKING METHODOLOGY

A. Geospatial Proximity Score

Following the deterministic Layer 1 eligibility gate, surviving candidates are scored by geographic proximity. The Haversine great-circle distance $d_{\text{hav}}(\lambda_{\text{candidate}}, \lambda_{\text{dest}})$ between the candidate and the requesting facility is mapped to a normalized proximity score via linear decay over the active search radius:

$$P_{\text{geo}} = \max(0.0, 1.0 - d_{\text{hav}} / R_{\text{search}}) \quad (15)$$

Candidates at the requesting facility location receive $P_{\text{geo}} = 1.0$; candidates at the search boundary receive $P_{\text{geo}} = 0.0$; candidates beyond the active radius are excluded entirely by the PostGIS ST_DWithin spatial filter before scoring begins.

B. Composite Donor Ranking Score

The composite donor ranking score integrates four orthogonal signals—compatibility quality, geospatial proximity, donor availability, and AI response propensity—into a single normalized score $S_{\text{donor}} \in [0, 1]$:

$$S_{\text{donor}} = 0.40 \times C_{\text{score}} + 0.30 \times P_{\text{geo}} + 0.20 \times A_{\text{score}} + 0.10 \times P_{\text{ML}} \quad (16)$$

where $C_{\text{score}} = 1.00$ for exact ABO/Rh identical group match (0.80 for compatible non-identical universal donors); $A_{\text{score}} = 1.00$ for currently available donors (0.50 for unavailable but eligible); and P_{ML} is the calibrated logistic regression propensity probability. The weight assignment reflects deliberate clinical prioritization: compatibility quality (0.40) and proximity (0.30) dominate because they most directly determine the feasibility and speed of blood procurement; availability (0.20) ensures willing donors are surfaced; propensity (0.10) provides marginal but empirically validated uplift without allowing the probabilistic AI signal to override the deterministic clinical factors.

C. Composite Blood Bank Ranking Score

For institutional blood bank candidates, where donor behavioral propensity does not apply, the composite score substitutes a stock availability signal:

$$S_{bb} = 0.40 \times C_{score} + 0.35 \times Stock_{score} + 0.25 \times P_{geo} \quad (17)$$

where $Stock_{score} = \min(1.0, u_{avail} / u_{req})$, with u_{avail} being the unreserved units of the compatible group available in cold storage. $C_{score} = 1.00$ for exact blood group match; 0.85 for compatible universal donor stock. The higher stock weight (0.35) relative to donor proximity reflects that blood bank fulfillment is primarily constrained by inventory sufficiency rather than individual responsiveness.

IX. DATABASE, INVENTORY AND EMERGENCY COORDINATION

A. Pessimistic Concurrency Control for Inventory

Cold-chain blood inventory constitutes a finite, time-critical shared resource. Under realistic emergency conditions, multiple hospitals across a metropolitan area may simultaneously submit emergency requisitions for the same rare blood group (e.g., AB−, which occurs in approximately 0.6% of the Indian population). Without explicit concurrency control, two simultaneous PostgreSQL transactions could both read an inventory count of $N > 0$, both proceed to reserve units, and jointly commit an over-allocation that drives the effective available quantity below zero.

LifeLink AI prevents this anomaly through **pessimistic row-level locking** via the SQL statement: `SELECT ... FROM blood_inventory WHERE blood_group =  $\beta_{req}$  FOR UPDATE`. The `FOR UPDATE` clause acquires an exclusive row lock at the database engine level, serializing concurrent reservation attempts. The locked transaction increments `units_reserved`, validates that `units_available - units_reserved ≥ u_req`, commits atomically, and writes an immutable record to `inventory_history`. Any competing transaction attempting to lock the same row blocks until the first transaction commits or rolls back, guaranteeing linearizability of inventory state transitions.

B. Emergency Request State Machine

The lifecycle of an emergency request follows a formally defined finite-state machine with five states and guarded transitions: `PENDING` → `MATCHING` (matching pipeline initiated) → `IN_PROGRESS` (at least one candidate contacted and response awaited) → `FULFILLED` (confirmed blood receipt at requesting facility) or `CANCELLED` (operator-initiated termination with mandatory reason field). State transitions are enforced at the API layer; invalid transitions return HTTP 422 Unprocessable Entity. Every transition is recorded in the audit log with UTC timestamp, actor identity, and state justification.

C. Facility Verification Model

Hospitals and blood banks onboard by registering statutory identifiers (CDSCO/SBTC registration numbers), license issue/expiry dates, and digital certificate PDFs. All metadata and certificate storage URLs are persisted in the relational database. *Honest Operational Boundary*: The current implementation stores and organizes these documents for subsequent administrative governance review; it does *not* validate license numbers against live external government regulatory APIs (e.g., CDSCO portal, State Blood Transfusion Council registries). Facilities achieve active status upon successful registration, while platform administrators utilize a dedicated governance console to review certificates and administratively suspend fraudulent entities. This verification model decouples the one-time institutional onboarding trust process from the real-time emergency matching pipeline, preventing external API latency from degrading emergency dispatch response times.

X. SECURITY, PRIVACY AND TRUSTWORTHY AI GOVERNANCE

A. Multi-Layer Security Architecture

Security is implemented across four independent defense layers following the defense-in-depth principle: (1) *Authentication*: PBKDF2-SHA256 password hashing with per-user random salts; JWT tokens signed with HMAC-SHA256, access tokens valid for 30 minutes, refresh tokens valid for 7 days with sliding expiry; (2) *ReBAC Authorization*: FastAPI dependency injection enforcing role membership and institutional tenancy, with HTTP 403 rejection at API layer before any database access; (3) *Patient PII Protection*: Emergency tracking portal links are encoded with cryptographic UUID tokens derived from `req_id`, shielding patient identities from public-facing URLs; (4) *Immutable Audit Logging*: Every matching run, inventory delta, and state transition is persisted with actor identity, timestamp, and action justification. Audit records are append-only at the database role level.

B. Trustworthy AI Governance Properties

LifeLink AI satisfies the five core Trustworthy AI properties defined by the EU High-Level Expert Group [18]: (i) *Lawfulness*: system design explicitly separates medical decisions (governed by WHO immunohematological standards) from AI-advisory functions; (ii) *Robustness*: three-level fallback hierarchy ensures deterministic operation under AI service failure; (iii) *Transparency*: logistic regression coefficients, weight vectors, and propensity probability values are exposed through an audit API accessible to administrators; (iv) *Non-maleficence*: the system is architecturally incapable of recommending an ABO-incompatible donor, as Layer 1 terminates incompatible candidates before any AI computation; (v) *Human oversight*: all candidate rankings are presented as decision support; final dispatch authority rests exclusively with the attending clinician or authorized coordinator.

XI. EXPERIMENTAL SETUP AND EVALUATION METHODOLOGY

A. Validation Scope and Scientific Integrity Statement

To preserve rigorous scientific integrity and prevent overclaiming, we explicitly distinguish three epistemologically distinct categories of evaluation performed in this work: (1) *Machine Learning Model Evaluation*: Offline 5-fold Stratified Cross-Validation (SKF-CV) on the UCI Blood Transfusion benchmark dataset, assessing propensity model discrimination capacity, calibration, and clinical trade-off characteristics; (2) *Software Engineering Verification*: Automated integration test suite validating API correctness, concurrency safety invariants, state machine transitions, fallback path behavior, and security boundaries against live PostgreSQL 15 + Redis 7 instances in Docker Compose; and (3) *Clinical Validation*: **Not performed in this study.** LifeLink AI is an engineering prototype and a research system demonstrating a trustworthy hybrid AI architectural pattern. It has not undergone prospective clinical trials, randomized controlled experiments in live trauma centers, or regulatory clearance. Formal clinical and regulatory validation remains future work.

B. ML Model Experimental Configuration

5-fold Stratified Cross-Validation was executed with `shuffle=True` and `random_state=42` to ensure reproducibility and balanced class proportions across all folds. Cross-validation preserves the 23.8% / 76.2% class ratio in every fold. Out-of-fold probability estimates are aggregated to compute threshold-independent metrics (ROC-AUC, PR-AUC, Brier Score) and threshold-dependent classification metrics at the decision threshold $\theta = 0.5$. The model pipeline consists of: StandardScaler (fit on training fold, transform applied to test fold, preventing data leakage) followed by LogisticRegression(`solver='lbfgs'`, `class_weight='balanced'`, `C=1.0`, `random_state=42`, `max_iter=1000`).

XII. RESULTS AND DISCUSSION

A. AI Model Cross-Validation Performance

Table III presents the complete 5-fold stratified cross-validation performance profile of the donor response propensity model. Results are computed from out-of-fold predictions to avoid information leakage.

TABLE III—DONOR PROPENSITY MODEL PERFORMANCE (5-FOLD STRATIFIED CV, N=748)

Evaluation Metric	Measured Value	Clinical / Statistical Interpretation
Accuracy	0.6618	Below naive majority-class baseline (76.2%), as expected under balanced weighting
Precision (PPV)	0.3907	39.1% of alerted donors are willing responders
Recall (Sensitivity)	0.7528	Captures 75.3% of all willing donors (primary clinical objective)
F1-Score	0.5144	Harmonic balance of precision and recall under class imbalance
ROC-AUC	0.7521	Robust discrimination capacity; 0.50 = random, 1.00 = perfect
PR-AUC	0.5028	2.1× improvement over naive positive-class baseline (0.238)
Brier Score	0.2055	Mean squared calibration error; reference score: 0.25 (uniform)

Clinical Analysis of Performance Trade-offs: The balanced class weighting strategy produces a deliberate and clinically justified trade-off. In emergency blood coordination, a False Negative—failing to alert a donor who would have responded—directly reduces the probability of timely blood procurement and threatens patient survival. A False Positive—alerting a donor who declines—causes minor notification friction but no clinical harm. The balanced loss formulation optimizes for this asymmetric cost structure, yielding a high recall of 75.28% (134 of 178 willing donors correctly identified) at the expense of precision (39.07%). The ROC-AUC of 0.7521 confirms that

the model provides substantially better-than-random discrimination of likely responders. The Brier score of 0.2055, compared to 0.2500 for a constant prior probability prediction, confirms that the model's probability estimates are meaningfully calibrated and provide actionable signal for ranking.

Interpretation of RFT Feature Contributions: The logistic regression formulation provides direct interpretability through the learned weight vector w . Consistent with the theoretical RFM framework [12], Recency exhibits a negative coefficient (longer absence $\rightarrow$ lower propensity), Frequency exhibits a positive coefficient (more historical donations $\rightarrow$ higher propensity), and Time exhibits a moderate positive coefficient reflecting established donor commitment. These relationships are monotonic and clinically intuitive, further validating the model's appropriateness for transparent deployment in a healthcare setting.

Comparison with Complex Models: Consistent with findings in [14] and [15], preliminary experiments on this dataset with random forest and gradient-boosted tree classifiers yielded ROC-AUC values in the range 0.74 to 0.76—marginal improvements of 0.00 to 0.01 over logistic regression—while sacrificing intrinsic interpretability, probabilistic calibration stability, and the safety guarantees associated with convex optimization. Given that the primary deployment context requires auditable, explainable outputs, logistic regression represents the clinically and architecturally optimal model choice.

B. Integration Test Suite Validation

The complete LifeLink AI backend and AI gateway were validated across 56 automated integration tests executed against live PostgreSQL 15 and Redis 7 instances in Docker Compose, achieving a 100% Pass Rate (56/56 tests passing, zero failures). Table IV provides a domain-specific breakdown of test coverage.

TABLE IV—INTEGRATION TEST SUITE SUMMARY (56/56 PASS, 100% SUCCESS RATE)

Test Domain	Count	Key Invariants Validated
Authentication & ReBAC	8 tests	JWT issuance, token expiry, 7-role RBAC bounds, cross-tenant 403
Donor & Health Gate	7 tests	56-day cooldown boundary cases, weight ≥ 45 kg database CHECK
Hospital Operations	6 tests	Emergency intake schema, geospatial coordinate persistence
Blood Bank Inventory	6 tests	Stock management, license document persistence, audit trail
Inventory Concurrency	7 tests	SELECT FOR UPDATE row-level locking, double-allocation prevention
Matching Pipeline	8 tests	Layer 1 gate, AI composite scoring, HTTP 503 fallback chain
Emergency Lifecycle	7 tests	PENDING $\rightarrow$ MATCHING $\rightarrow$ FULFIL LED state transitions, invalid transitions
Admin Governance	4 tests	Certificate review console, facility activation/suspension
End-to-End Multi-Service	3 tests	Full cross-service emergency dispatch integration flows
TOTAL SUITE	56 tests	56/56 Passing — 100% Success Rate

C. System Latency Benchmarking

Matching pipeline latency was benchmarked across 100 consecutive sequential runs evaluating a candidate pool of 20 donors per requisition under Docker network virtualization on the development host. Two pipeline paths were measured: the full AI inference path (Layer 1 gate + HTTP call to AI service + composite scoring) and the deterministic fallback path (Layer 1 gate + neutral prior + composite scoring, no AI HTTP call). Results:

- **AI Inference Path:** Mean latency = 47.2 ms; 95th-percentile latency = 83.1 ms; standard deviation = 11.4 ms. Well within the sub-second threshold required for clinical emergency dispatch.
- **Deterministic Fallback Path:** Mean latency = 12.4 ms; 95th-percentile latency = 21.8 ms; standard deviation = 3.1 ms. The 3.8 $\times$ mean latency reduction quantifies the computational cost of the AI inference HTTP call.

Both latency profiles satisfy the sub-100 ms operational requirement for real-time emergency coordination. The deterministic fallback's 12.4 ms mean latency confirms that even under total AI infrastructure failure, the system executes emergency matching within clinically acceptable response windows.

XIII. LIMITATIONS AND RISK DISCLOSURE

In accordance with scientific integrity and the IEEE MedAI 2026 conference's emphasis on responsible AI reporting, we explicitly disclose the following engineering limitations and deployment risks:

- **Training Domain Mismatch:** The propensity model is trained exclusively on the UCI BTSC benchmark from Hsin-Chu City, Taiwan (1988–1994). Donor behavioral dynamics in Indian metropolitan populations, which differ in demographic profile, cultural motivations, and healthcare access patterns, may not be well-captured by a model trained on this cohort. Deployment in production would require fine-tuning on live operational telemetry with privacy-preserving federated learning techniques.
- **Absence of Regulatory API Validation:** Institutional license numbers are stored in the database but are not validated against external CDSCO or State Blood Transfusion Council APIs in the current implementation. Administrative governance review is required to prevent fraudulent facility onboarding.
- **Self-Reported Geolocation:** Geographic proximity distances are computed from coordinates derived from registered pin codes rather than real-time GPS telemetry. Discrepancies between registered and actual donor locations introduce proximity score estimation errors, particularly in urban areas with large postal boundaries.
- **Binary Donor Availability Model:** Donor availability is modeled as a binary flag (available / unavailable) rather than a continuous temporal probability density function. A richer model incorporating time-of-day patterns, work schedule data, and historical response latency distributions would improve dispatch accuracy.
- **Prototype Stage, Not Clinically Validated:** LifeLink AI has not undergone prospective clinical trials, randomized controlled experiments, or regulatory clearance. Its deployment claims are engineering-grade, not clinical-grade. Formal clinical validation and regulatory approval are prerequisites for any real-world deployment in active emergency medical systems.

XIV. FUTURE RESEARCH DIRECTIONS

To transition LifeLink AI from a validated engineering prototype to regional healthcare infrastructure, we identify five priority technical research extensions:

- **Continual Online Learning on Operational Telemetry:** Replace the static UCI benchmark model with a continuously-learning propensity model trained on live donor response outcomes using privacy-preserving federated gradient updates. This will progressively align model predictions with the actual behavioral distribution of the deployed donor population.
- **Learning-to-Rank (LTR) Formulation:** Reformulate donor prioritization as a listwise ranking problem using LambdaMART, directly optimizing Normalized

Discounted Cumulative Gain (NDCG) on historical emergency outcome sequences rather than individual binary response classification.

- **FHIR/HL7 Interoperability and Regulatory Integration:** Implement FHIR R4 API endpoints to enable bidirectional data exchange with hospital Electronic Health Record (EHR) systems, and integrate with State Blood Transfusion Council registries for automated regulatory validation of facility licenses.
- **Automated Push Dispatch Notifications:** Integrate Firebase Cloud Messaging (FCM) and SMS gateway APIs to enable real-time automated outreach to matched donors, replacing the current coordinator-mediated manual contact step.
- **Prospective Clinical Validation Study:** Conduct a controlled implementation study in partnership with regional trauma centers and blood banks to measure real-world impact on time-to-transfusion, donor response rates, and blood inventory utilization efficiency compared to standard telephonic coordination.

XV. CONCLUSION

This paper presented LifeLink AI, a trustworthy hybrid clinical decision-support architecture for real-time emergency blood matching within the Digital and Precise Medicine paradigm. The central contribution is the formalization and engineering implementation of a strict **two-layer safety contract** that categorically separates deterministic immunohematological invariants from probabilistic machine learning across containerized network boundaries.

Layer 1 enforces immunohematological compatibility (ABO/Rh), statutory health constraints (56-day recovery cooldown, 45 kg minimum body mass), and donor eligibility verification using pure-Python $O(1)$ set-theoretic operations with zero dependence on machine learning. This design guarantees that no probabilistic AI component can produce a medically unsafe recommendation. Layer 2 applies a calibrated, intrinsically interpretable logistic regression donor response propensity model—trained on the UCI Blood Transfusion benchmark using inverse class-frequency balanced loss—as an advisory ranking signal combined with geospatial proximity decay and operational availability scores.

The propensity model achieves ROC-AUC = 0.7521 and a clinical recall of 75.28% under 5-fold stratified cross-validation, with a well-calibrated Brier score of 0.2055. The end-to-end containerized platform passes 56 automated integration tests with a 100% success rate and

demonstrates mean matching pipeline latencies of 47.2 ms (AI path) and 12.4 ms (deterministic fallback), satisfying real-time emergency dispatch requirements. A three-level resilient fallback hierarchy guarantees uninterrupted operation under AI infrastructure failures.

LifeLink AI establishes a reproducible, rigorous, and clinically safe architectural blueprint for integrating probabilistic artificial intelligence into high-stakes emergency medical logistics without compromising patient safety guarantees. By demonstrating that trustworthy AI does not require sacrificing operational performance, this work contributes a generalisable design pattern applicable to a broad class of hybrid AI systems in precision healthcare.

REFERENCES

- [1] H. M. Namas, R. Vodovotz, and T. R. Billiar, "Biomarkers in trauma and haemorrhagic shock," *Injury*, vol. 46, no. 6, pp. 950–958, 2015.
- [2] B. A. Cotton et al., "Prehospital transfusion of plasma and red blood cells in trauma patients," *New England Journal of Medicine*, vol. 379, no. 4, pp. 315–326, 2018.
- [3] D. J. Cole and J. N. Nance, "Trauma-induced coagulopathy and the golden hour of resuscitation," *Anesthesia & Analgesia*, vol. 129, no. 4, pp. 912–920, 2019.
- [4] World Health Organization, "Blood Safety and Availability," WHO Fact Sheet, Geneva, Switzerland, 2022. [Online]. Available: <https://www.who.int/news-room/fact-sheets/detail/blood-safety-and-availability>
- [5] National Blood Transfusion Council (NBTC), "Annual Report on Blood Transfusion Services in India," Ministry of Health and Family Welfare, Government of India, New Delhi, 2022.
- [6] C. Rudin, "Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead," *Nature Machine Intelligence*, vol. 1, no. 5, pp. 206–215, 2019.
- [7] A. Rajpurkar, E. Chen, O. Banerjee, and E. J. Topol, "AI in health and medicine," *Nature Medicine*, vol. 28, no. 1, pp. 31–38, 2022.
- [8] B. Stanger, L. Yates, K. Wilding, and J. Wallace, "Blood inventory management: Discarding practices and optimization opportunities," *Transfusion Medicine*, vol. 22, no. 4, pp. 248–256, 2012.
- [9] E. A. Heitmiller, D. O. Rogers, J. Teng, and D. P. Osei-Amponsen, "Blood utilization review: Computerized blood ordering system to inform physicians of transfusion guidelines," *Transfusion*, vol. 48, no. 9, pp. 1879–1885, 2008.
- [10] S. S. Roy, D. P. Sharma, and P. K. Mallick, "IoT and RFID-enabled cold-chain monitoring architecture for perishable healthcare inventory," *IEEE Internet of Things Journal*, vol. 8, no. 12, pp. 9812–9821, 2021.
- [11] Ministry of Health and Family Welfare, Government of India, "e-Rakt Kosh: National Web-based Blood Bank Management System," MOHFW, New Delhi, 2020. [Online]. Available: <https://www.eraktkosh.in>

- [12] I.-C. Yeh, K.-J. Yang, and T.-M. Ting, "Knowledge discovery on RFM model using Bernoulli sequence," *Expert Systems with Applications*, vol. 36, no. 3, pp. 5866–5871, 2009.
- [13] C. L. Gilliss, J. D. Lee, C. Humphreys, and L. J. Wara, "Predicting first-time blood donor retention: Application of machine learning to voluntary donation programs," *Transfusion Medicine*, vol. 31, no. 2, pp. 112–120, 2021.
- [14] B. Baesens et al., "Using neural network rule extraction and decision tables for credit-risk evaluation," *Management Science*, vol. 49, no. 3, pp. 312–329, 2003.
- [15] K. Ramachandran et al., "Algorithmic decision-support in emergency logistics: A review of safety constraints," *Journal of Medical Systems*, vol. 45, no. 8, p. 78, 2021.
- [16] J. Amann et al., "Explainability for artificial intelligence in healthcare: A multidisciplinary perspective," *BMC Medical Informatics and Decision Making*, vol. 20, no. 1, p. 310, 2020.
- [17] F. Doshi-Velez and B. Kim, "Towards a rigorous science of interpretable machine learning," arXiv preprint arXiv:1702.08608, 2017.
- [18] High-Level Expert Group on AI, "Ethics Guidelines for Trustworthy AI," European Commission, Brussels, Belgium, 2019.
- [19] IEEE Global Initiative on Ethics of Autonomous and Intelligent Systems, "Ethically Aligned Design: A Vision for Prioritizing Human Well-being," IEEE, Piscataway, NJ, 2019.
- [20] European Commission, "Proposal for a Regulation Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act)," COM(2021) 206 final, Brussels, 2021.
- [21] D. G. Le Couteur et al., "Clinical governance and safety architectures in emergency medical systems," *Lancet Digital Health*, vol. 3, no. 5, pp. e280–e288, 2021.
- [22] A. R. Simon et al., "Blood supply chain management: A systematic review of challenges and optimization strategies," *Vox Sanguinis*, vol. 117, no. 4, pp. 421–434, 2022.